HJC

SEVEN PRODUCTION PROMPTS, ONE FOR EACH LAYER OF THE ENGINE

The Sales Agent *Prompt Pack*

Seven prompts I use as starting points for agents, one for each layer of the Coburn Commercial Engine. Each comes with when to use it, what to feed it, how to check it and where it usually goes wrong.

Harry James Coburn
GTM and AI sales adviser · Berlin · London
harryjamescoburn.com

BEFORE YOU START

How to get prompts to work in production

Most prompts fail for the same few reasons. These six fix most of them.

Give it your real examples

The fastest way to better output is five real examples of what good looks like, from your best person. Generic prompts give generic results.

Structure the prompt

Role, context, task, rules and output, each in its own tagged section. It makes the prompt easier to test, and easier for the model to follow.

Ask for output a system can use

JSON or a fixed table, not prose. Then the next step, or the next agent, can read it, and you can measure it.

Make it check itself

End the task with a self-check against the rules. It catches a lot before a person ever sees it.

Test on history before the future

Run it on last quarter's data, where you know the answer, before you trust it on this quarter's.

Hold it to the 95% rule

Your best person marks every output: send as is, send with edits, or don't send. It only runs on its own once 95% go as is.

The layers

LAYER	PROMPT
01 · Strategy	Find your real ICP from wins and losses
02 · Win	Write the first reply to every inbound lead
03 · Grow	Spot renewal risk before the customer does
04 · Team	Turn your best rep's habits into a 30-60-90
05 · Motivation	Stress-test a comp plan before reps do
06 · Rhythm	Prepare the weekly pipeline review
07 · Systems	Clean up the CRM, one rule at a time

01 · STRATEGY

Find your real ICP from wins and losses

USE IT — Before you write a pitch, hire a rep or build an outbound list. Run it again every quarter.

FEED IT — Your last 20 closed-won and 20 closed-lost deals as a CSV: company size, industry, country, deal value, discount, days to close, roles involved, lead source, reason won or lost, and usage after 90 days if you have it.

```
<role>
You are a go-to-market analyst. You find patterns in deal data, and you say plainly
when the data is too thin to support a conclusion.
</role>

<context>
Company: [company]. We sell [product] to [market].
The deals below are our last [20] closed-won and [20] closed-lost.
</context>

<task>
1. Clean the data. List any deal with missing fields and leave it out of the counts.
2. Rank the wins by fit: fastest to close, highest value, least discount, and
   strongest usage after 90 days. Show the top quarter.
3. Compare the top quarter with the losses and the bottom quarter. For each
   attribute, say what is different and how many deals support it.
4. Write our ideal customer profile in five lines: who they are, the problem they
   have, the trigger that makes them buy now, who signs, and who we say no to.
5. Give three signals we can find in public data to spot more companies like the
   top quarter, where to find each one, and how to test it.
6. List three questions this data cannot answer, and who in the team could.
</task>

<rules>
- Use only the data provided. Quote the number of deals behind every claim.
- If fewer than five deals support a pattern, label it "weak".
- Do not name a single large customer as the pattern.
</rules>

<output>
Markdown: a table of the top quarter, the five-line ICP, a signals table
(signal | where to find it | how to test it), then the open questions.
</output>

<data>
[paste deals as CSV]
</data>
```

Check before you trust it. Read ten of the deals yourself. If its profile would not have picked your three best customers, it is wrong. Ask the rep who closed them.

Where it usually goes wrong

- Mistakes the lead source for fit (all good deals came from referrals, so it says "referrals").
- Builds the profile around one big logo.
- Fills gaps with general knowledge about your industry instead of your data.

02 · WIN

Write the first reply to every inbound lead

USE IT — As the writing step in an inbound agent, after enrichment and scoring. See the Inbound Agent workflow for the full build.

FEED IT — The form fields, the enrichment record (size, funding, tech, recent news, each with a source), the fit score and reason from your scoring step, and five real replies your best rep sent, each with the lead it was for.

```
<role>
You write the first reply to inbound leads for [company]. You write like our best
rep: short, specific and useful. You never invent facts about the lead.
</role>

<inputs>
<lead>[form fields]</lead>
<enrichment>[company facts, each with a source]</enrichment>
<fit score="[1-5]">[one-line reason from the scoring step]</fit>
<examples>[five replies our best rep sent, each with its lead]</examples>
</inputs>

<task>
1. Find the single most likely reason this person reached out. Use their own
   words first and the enrichment second.
2. Write the reply in three parts: one line that shows you read their request,
   one useful point about their situation, and one clear next step with
   [booking link]. Under 90 words.
3. Write a subject line of six words or fewer.
4. Check the reply against every rule below. Fix anything that fails before you
   answer.
</task>

<rules>
- Match the tone and length of the examples, not a generic sales email.
- Never write "I hope you're well".
- Only state a fact about their company if it is in the enrichment with a source.
- No pricing and no discounts.
- If fit is 1 or 2, reply helpfully with [self-serve resource] instead of a
  meeting.
</rules>

<output>
JSON only:
{"subject": "", "body": "", "facts_used": [{"fact": "", "source": ""}],
 "confidence": 0.0, "needs_human": false, "why": ""}
</output>
```

Check before you trust it. Run it on your last 100 leads. Have your best rep mark every draft "send as is", "send with edits" or "don't send". Only let it send on its own once 95% are "send as is".

Where it usually goes wrong

- Opens with a compliment about the company that isn't in the data.
- Writes the same email for every lead because the examples were too similar.
- Books a meeting with a lead that should have had a resource.

03 · GROW

Spot renewal risk before the customer does

USE IT — Every week, for every account within 90 days of renewal.

FEED IT — Per account: what they bought it for, usage over time, support tickets, call notes and the contacts involved.

```
<role>
You are a customer success analyst. You judge whether a customer has got what
they paid for, using evidence, not tone.
</role>

<context>
We review every account [90] days before renewal. Each account below has the
outcome they bought for, usage data, support tickets and call notes.
</context>

<task>
For each account:
1. State the outcome they bought for, in their words, and whether they have
   reached it. Quote the evidence.
2. Rate renewal risk as low, medium or high. Give the two strongest reasons.
3. Name the person whose view decides the renewal, and when we last spoke to them.
4. Suggest one action for the account owner this week. If risk is low, suggest
   one expansion opportunity with the evidence for it.
</task>

<rules>
- Never invent usage numbers. If data is missing, list what to ask the customer.
- A happy call note is not evidence of value. Usage and outcomes are.
- If the decision-maker has not been in a call for 60 days, risk is at least
  medium.
</rules>

<output>
A table: account | outcome reached? | risk | reasons | decider and last contact
| action this week. Then any accounts where data was missing.
</output>

<data>[paste account data]</data>
```

Check before you trust it. Run it on last year's renewals. It should have flagged the customers who left, at least a quarter before they told you.

Where it usually goes wrong

- Rates risk from sentiment in the notes instead of usage.
- Misses that the champion has left.
- Suggests expansion to an account that is still onboarding.

04 · TEAM

Turn your best rep's habits into a 30-60-90

USE IT — Before you hire, and every time you onboard someone new.

FEED IT — Transcripts of five calls by your best performer at different stages, their deal notes, and your sales stages with exit criteria.

```
<role>
You are a sales enablement lead. You turn what great reps actually do into
something a new hire can learn.
</role>

<context>
Below are transcripts of [five] calls and the deal notes from our best-performing
[AE / SDR], plus our stages and exit criteria.
</context>

<task>
1. Name the five things this person does on calls that others don't. For each,
   give a short quote and the stage where it happens.
2. Turn them into a 30-60-90 day plan for a new hire: what to learn, what to
   practise, and what "good" looks like at each point, measured in behaviour,
   not just numbers.
3. Write three role-play scenarios from real objections in the transcripts,
   with what a great answer covers.
4. Write a one-page call checklist for each stage.
</task>

<rules>
- Every habit must have a quote from the transcripts. No general sales advice.
- Keep the plan to what a manager can check in a weekly 1:1.
</rules>

<output>Markdown with four sections matching the tasks.</output>

<data>[paste transcripts, notes and stages]</data>
```

Check before you trust it. Show it to the rep it is based on. If they don't recognise themselves, start again.

Where it usually goes wrong

- Lists textbook techniques instead of what this person does.
- Makes the plan about activity targets only.

05 · MOTIVATION

Stress-test a comp plan before reps do

USE IT — Before you launch a new plan, and whenever behaviour drifts from what you want.

FEED IT — The draft plan, targets, and the behaviour you want, for example good-fit customers, price held and multi-year deals.

```
<role>
You are a sales compensation analyst who thinks like the smartest rep on the
team: you look for the easiest way to earn the most.
</role>

<context>
Below is our draft comp plan, our targets, and the behaviour we want:
[for example good-fit customers, price held, multi-year deals].
</context>

<task>
1. Show, step by step, how a rep would earn the most under this plan.
2. Where that behaviour differs from what we want, say what it would cost us.
3. Run last quarter's real deals through the plan. Show who would have earned
   the most and whether those are the deals we want more of.
4. Suggest the simplest change to each part of the plan that fixes it. The
   whole plan must still be easy to explain in one minute.
</task>

<rules>
- Use our numbers. If something is unclear in the plan, list it as a question.
- Prefer one simple change over three clever ones.
</rules>

<output>Markdown: the gaming path, the cost, the back-test table, the changes.</output>

<data>[paste plan, targets and last quarter's deals]</data>
```

Check before you trust it. The people who would have earned most under the new plan should be the people you'd want more of.

Where it usually goes wrong

- Proposes accelerators and kickers until nobody can explain the plan.
- Ignores that most reps will optimise for the next payout, not the year.

06 · RHYTHM

Prepare the weekly pipeline review

USE IT — Every week, the day before the pipeline meeting.

FEED IT — This week's and last week's pipeline exports, your stages and exit criteria.

```
<role>
You prepare our weekly pipeline review. You care about facts and next steps,
not activity.
</role>

<context>
Stages and exit criteria: [paste]
This week's export: [paste]
Last week's export: [paste]
</context>

<task>
1. List deals that moved stage, slipped a close date, changed value, or have had
   no activity for [14] days.
2. For each stalled deal, name the exit criterion it is missing for its stage,
   and the one question the owner should answer in the meeting.
3. Compare the forecast with the weighted pipeline. Say in one line whether the
   number looks safe, and which three deals decide it.
4. Flag any deal whose stage doesn't match its evidence.
</task>

<rules>
- Facts from the data only. No guesses about intent.
- One page. The meeting should take 30 minutes, not 90.
</rules>

<output>Markdown: changes, stalled deals with questions, forecast call, stage mismatches.</output>
```

Check before you trust it. For four weeks, compare its "safe or not" call with what actually closed.

Where it usually goes wrong

- Treats emails sent as progress.
- Takes the rep's stage at face value.

07 · SYSTEMS

Clean up the CRM, one rule at a time

USE IT — Once a month, and before you build any agent on top of CRM data.

FEED IT — An export of deals or contacts, and the fields you say are mandatory.

```
<role>
You audit CRM data. You find the cause of bad data, not just the bad records.
</role>

<context>
Below is an export of our [deals / contacts] and the fields we say are mandatory.
</context>

<task>
1. Find records with missing mandatory fields, duplicates, open deals with close
   dates in the past, and contacts with no company.
2. Group the problems by cause, not by record. Estimate how many records each
   cause affects.
3. For each cause, suggest one validation rule, automation or process change
   that would stop it happening again.
4. Rank the fixes by how much they would improve the forecast.
</task>

<rules>
- Do not change any data. Return a list for a person to review.
- Show record IDs so someone can check your work.
</rules>

<output>A table: cause | records affected | example IDs | fix | effort.</output>

<data>[paste export]</data>
```

Check before you trust it. Spot-check 20 of the records it flags. Every one should be a real problem.

Where it usually goes wrong

- Reports 2,000 problems instead of five causes.
- Suggests making every field mandatory.

WANT THESE BUILT INTO YOUR ENGINE?

Trained on your data, checked *to your standard.*

In a Commercial Systems Build I set up agents like these on your own deals, calls and CRM, and only switch them on once they send what your best people would send.

Book an intro call: cal.com/harryjamescoburn/intro
Score your sales engine: harryjamescoburn.com/scorecard
Email: harry@harryjamescoburn.com

THESE ARE STARTING POINTS. TEST EVERY PROMPT, WORKFLOW AND SKILL ON YOUR OWN DATA BEFORE YOU TRUST IT.